



# NEPTUNE

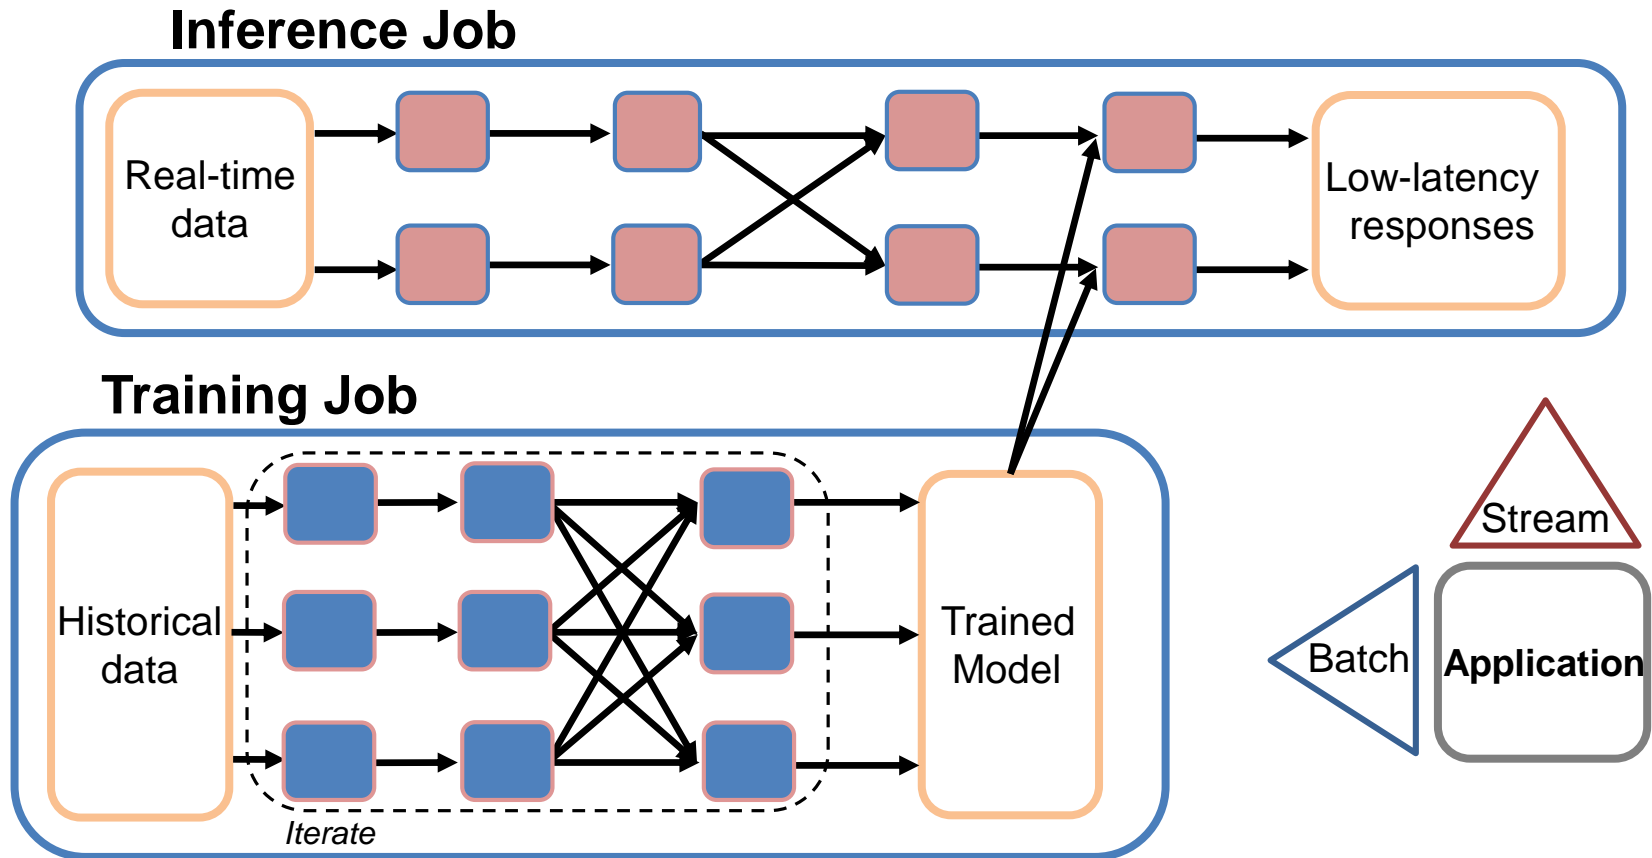
## Scheduling Suspendable Tasks for Unified Stream/Batch Applications

**Panagiotis Garefalakis**  
Imperial College London  
[pgaref@imperial.ac.uk](mailto:pgaref@imperial.ac.uk)

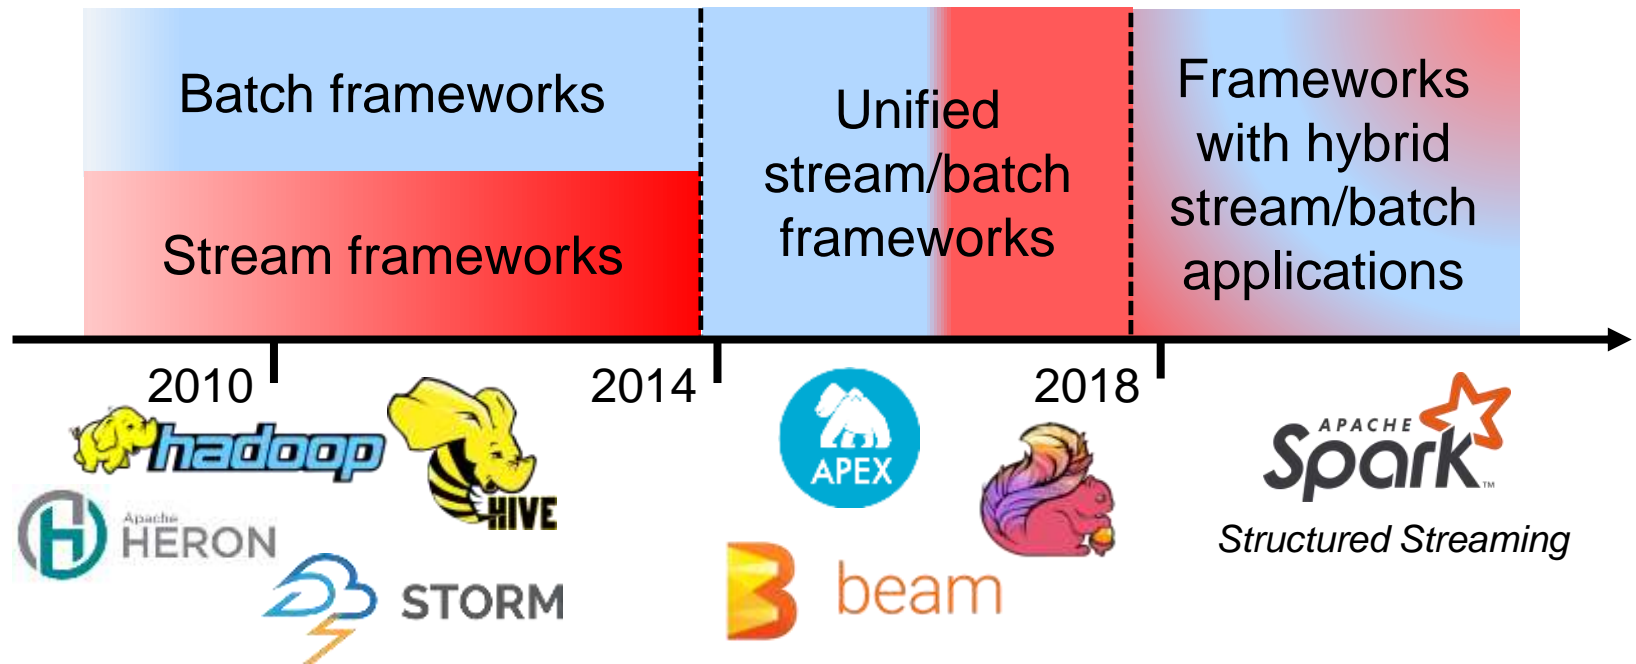
**Konstantinos Karanasos**  
Microsoft  
[kokarana@microsoft.com](mailto:kokarana@microsoft.com)

**Peter Pietzuch**  
Imperial College London  
[prp@imperial.ac.uk](mailto:prp@imperial.ac.uk)

# Unified application example



# Evolution of analytics frameworks



# Stream/Batch application requirements

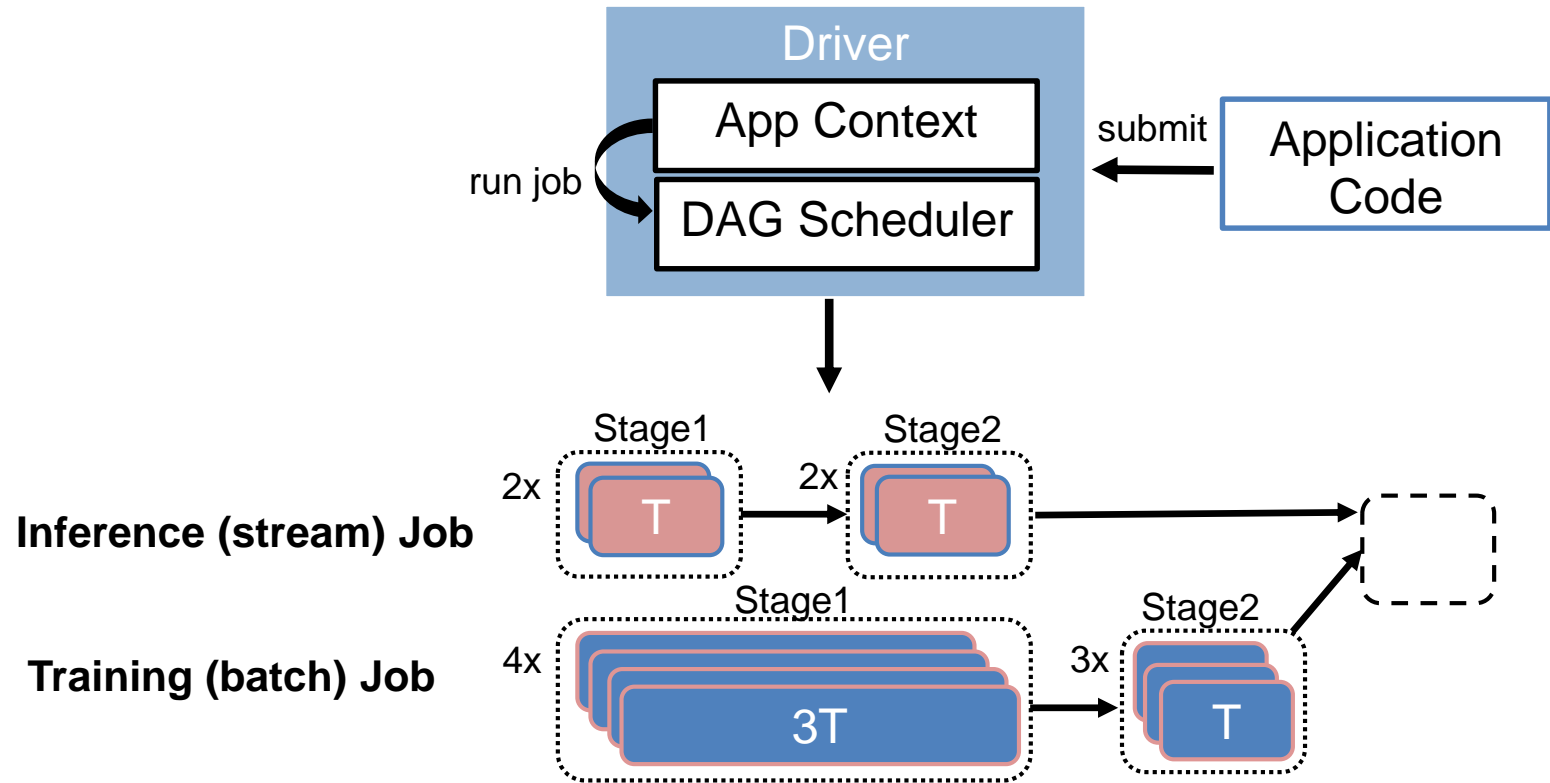
## Requirements

- > **Latency:** Execute inference job with minimum delay
- > **Throughput:** Batch jobs should not be compromised
- > **Efficiency:** Achieve high cluster resource utilization

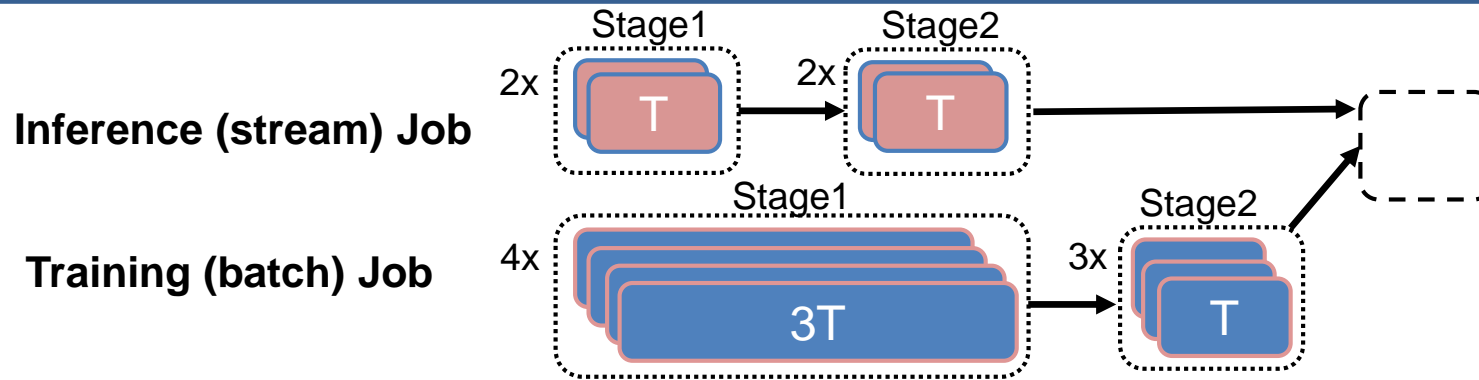


Challenge: schedule stream/batch jobs to satisfy their diverse requirements


# Stream/Batch application scheduling



# Stream/Batch application scheduling

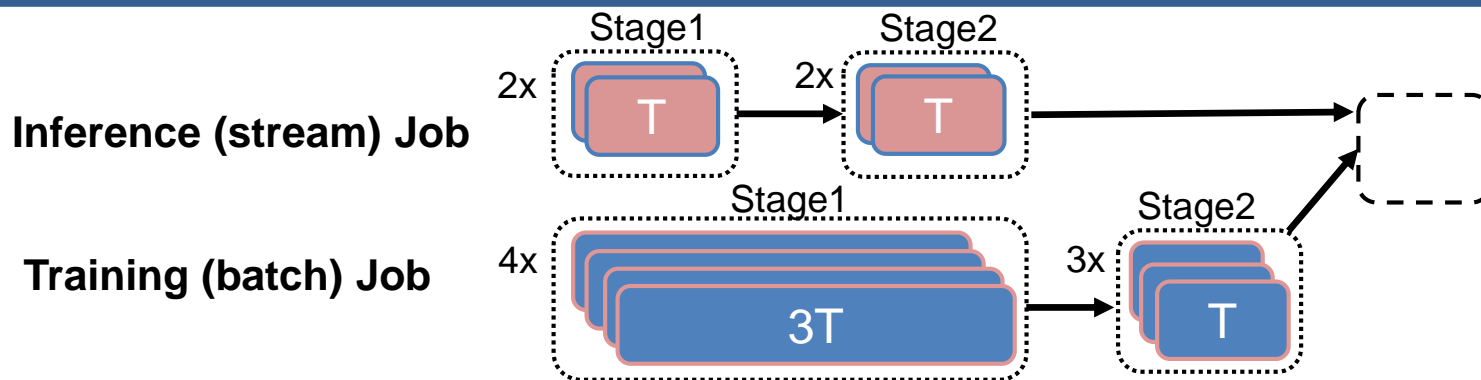


> **Static allocation:** dedicate resources to each job




Resources can not be shared across jobs

# Stream/Batch application scheduling

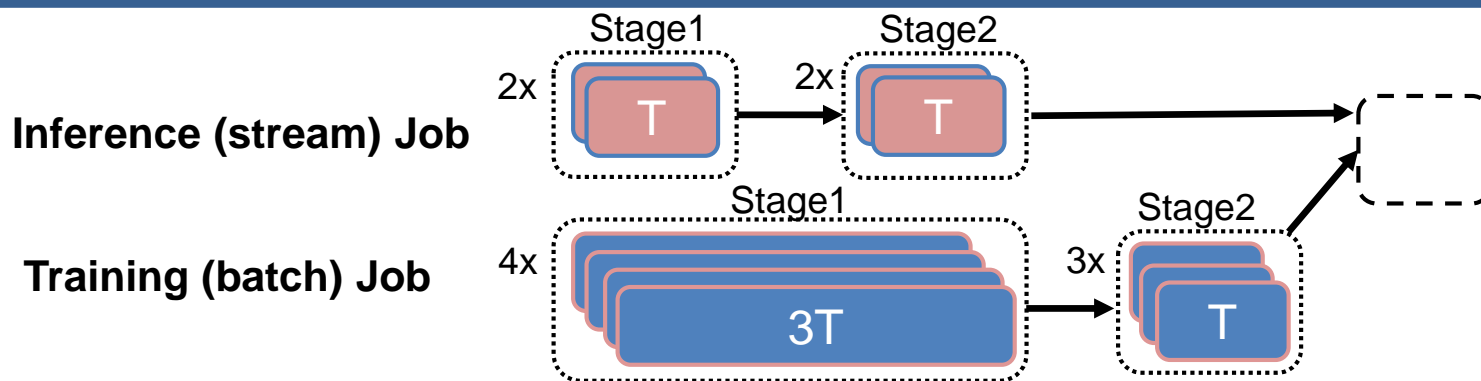


> **FIFO:** first job runs to completion




Long batch jobs increase stream job latency

# Stream/Batch application scheduling




> **FAIR:** weight share resources across jobs




Better packing with non-optimal latency



# Stream/Batch application scheduling

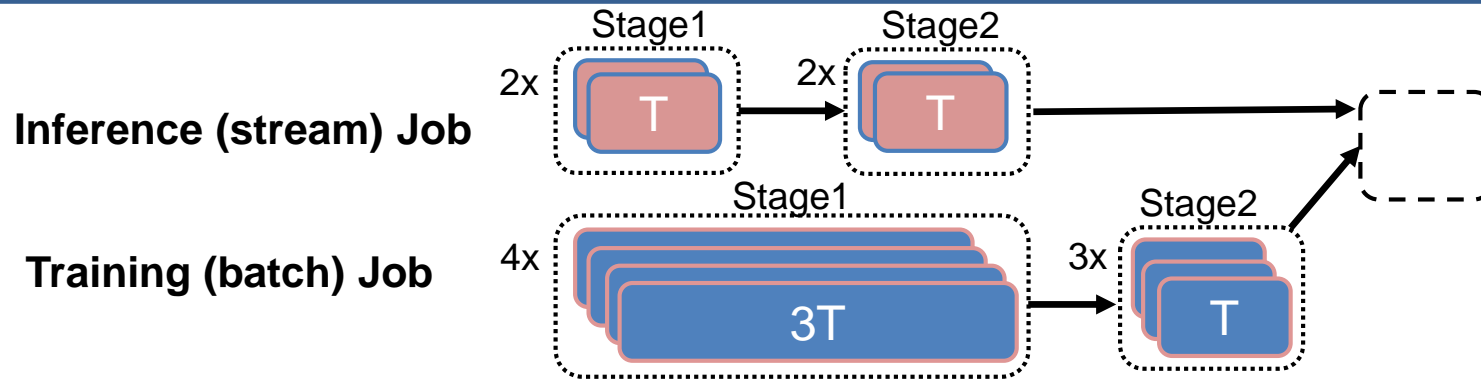


> **KILL:** avoid queueing by preempting batch tasks




Better latency at the expense of extra work

# Stream/Batch application scheduling



> **NEPTUNE:** minimize queueing and wasted work!



# Challenges

---

- > How to **minimize queuing** for latency-sensitive jobs and wasted work?
- > How to **natively** support stream/batch applications?
- > How to **satisfy** different stream/batch application **requirements** and high-level objectives?

# NEPTUNE

Execution framework for Stream/Batch applications

- > How to **minimize queuing** for latency-sensitive jobs and wasted work?



Support suspendable tasks

- > How to **natively** support stream/batch applications?



Unified execution framework on top of



*Structured Streaming*


- > How to **satisfy** different stream/batch application requirements and high-level objectives?



Introduce pluggable scheduling policies

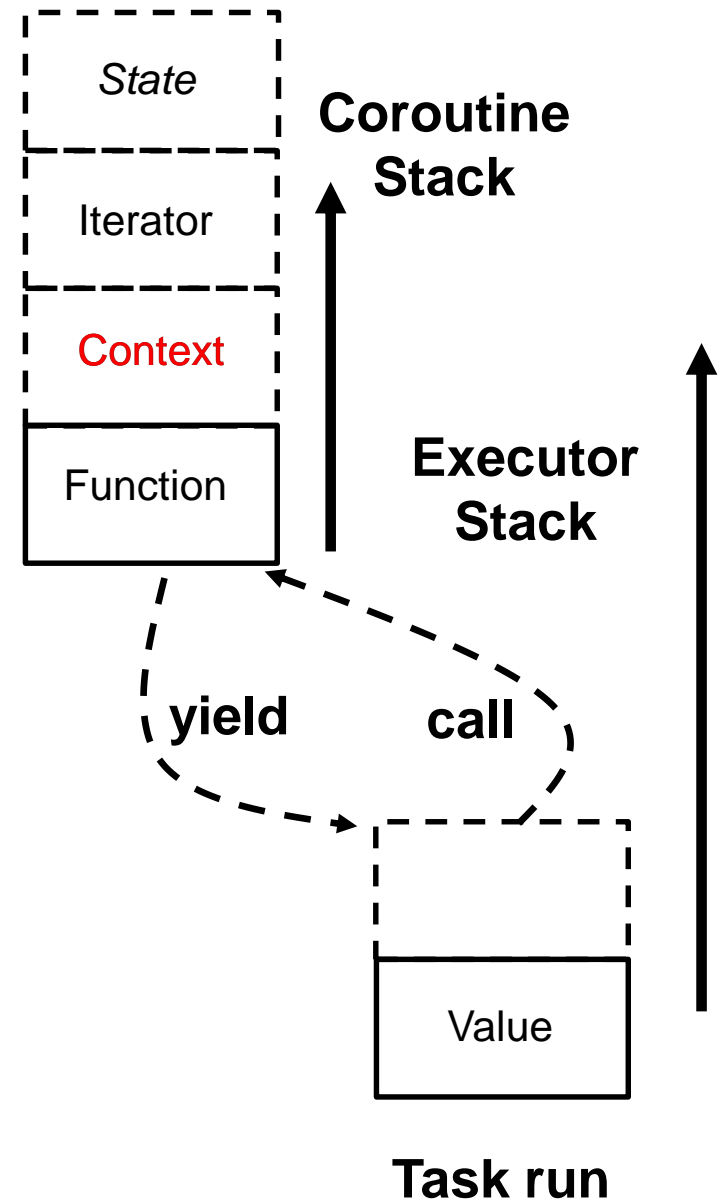
# Typical tasks

- > **Tasks:** apply a function to a partition of data
- > Subroutines that run in executor to completion
- > **Preemption problem:**
  - > Loss of progress (kill)
  - > Unpredictable preemption times (checkpointing)



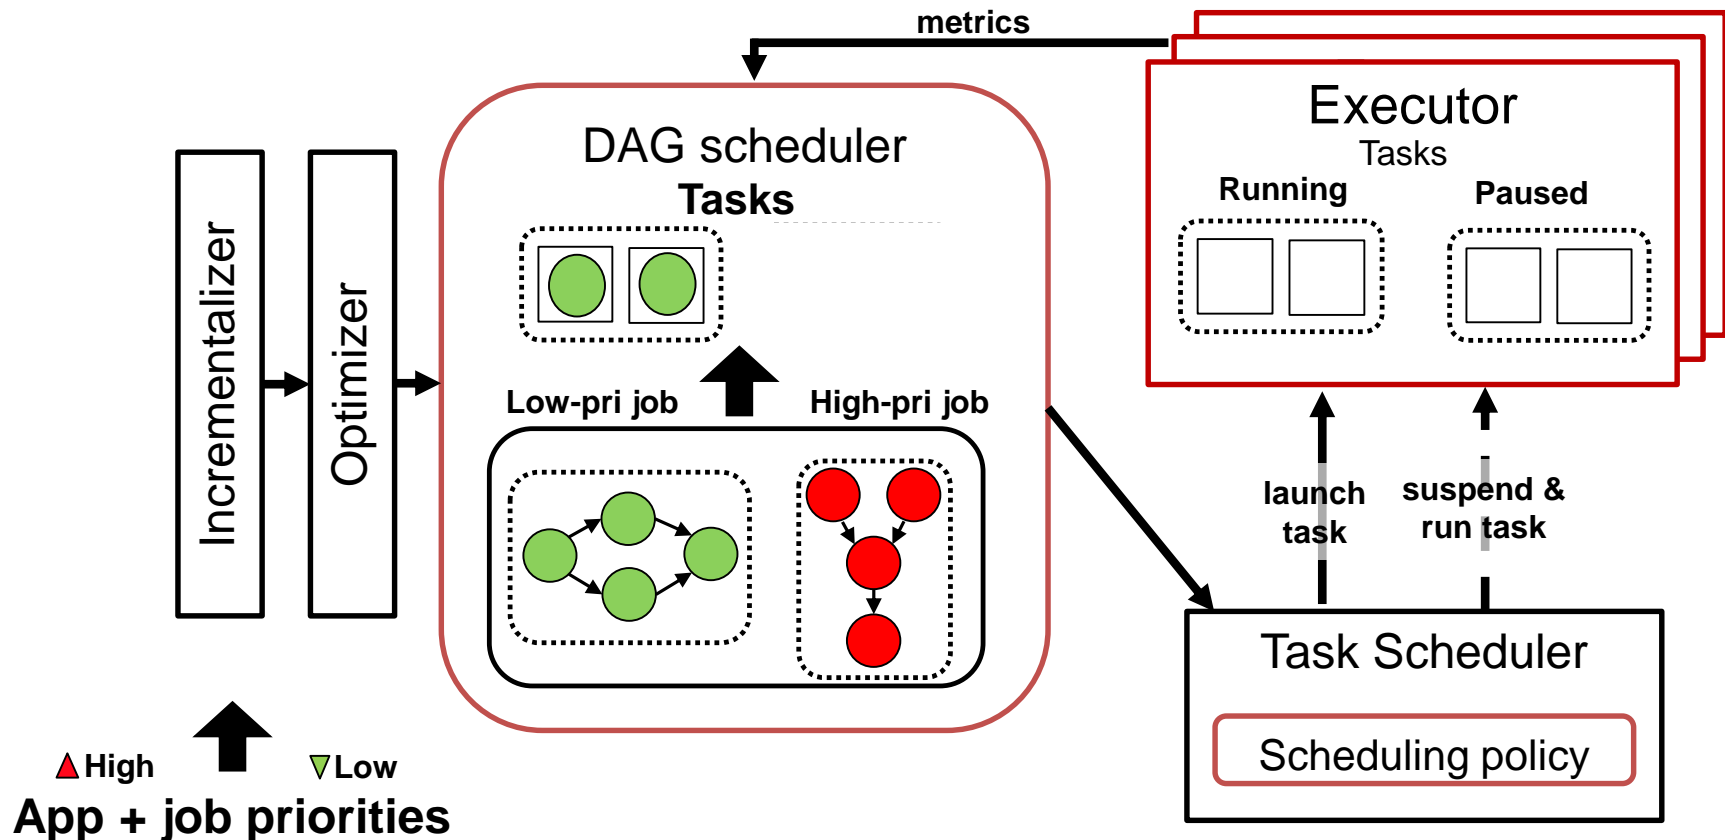
# Suspendable tasks

- > **Idea:** use coroutines
  - > Separate stacks to store task state
  - > **Yield** points handing over control to the executor
- > **Cooperative preemption:**
  - > Suspend and resume in *milliseconds*
  - > Work-preserving
- > Transparent to the user



# Execution framework

- > **Problem:** not just assign but also suspend and resume
- > **Idea:** centralized scheduler with pluggable policies




# Scheduling policies

- > **Idea:** policies trigger task suspension and resumption
  - > Guarantee that stream tasks bypass batch tasks
  - > Satisfy higher-level objectives i.e. balance cluster load
  - > Avoid **starvation** by suspending up to a number of times
- > **Load-balancing** (LB): takes into account executors' memory conditions and equalize the number of tasks per node
- > **Locality- and memory aware** (LMA): respect task locality preferences in addition to load-balancing



# Implementation

- > Built as an extension to  **2.4.0** (<https://github.com/lsds/Neptune>)
- > Ported all ResultTask, ShuffleMapTask functionality across programming interfaces to coroutines
- > Extended Spark's DAG Scheduler to allow job stages with different requirements (priorities)
- > Added additional Executor performance metrics as part of the heartbeat mechanism

# Azure deployment

---


## > Cluster

- 75 nodes with 4 cores and 32 GB of memory each

## > Workloads

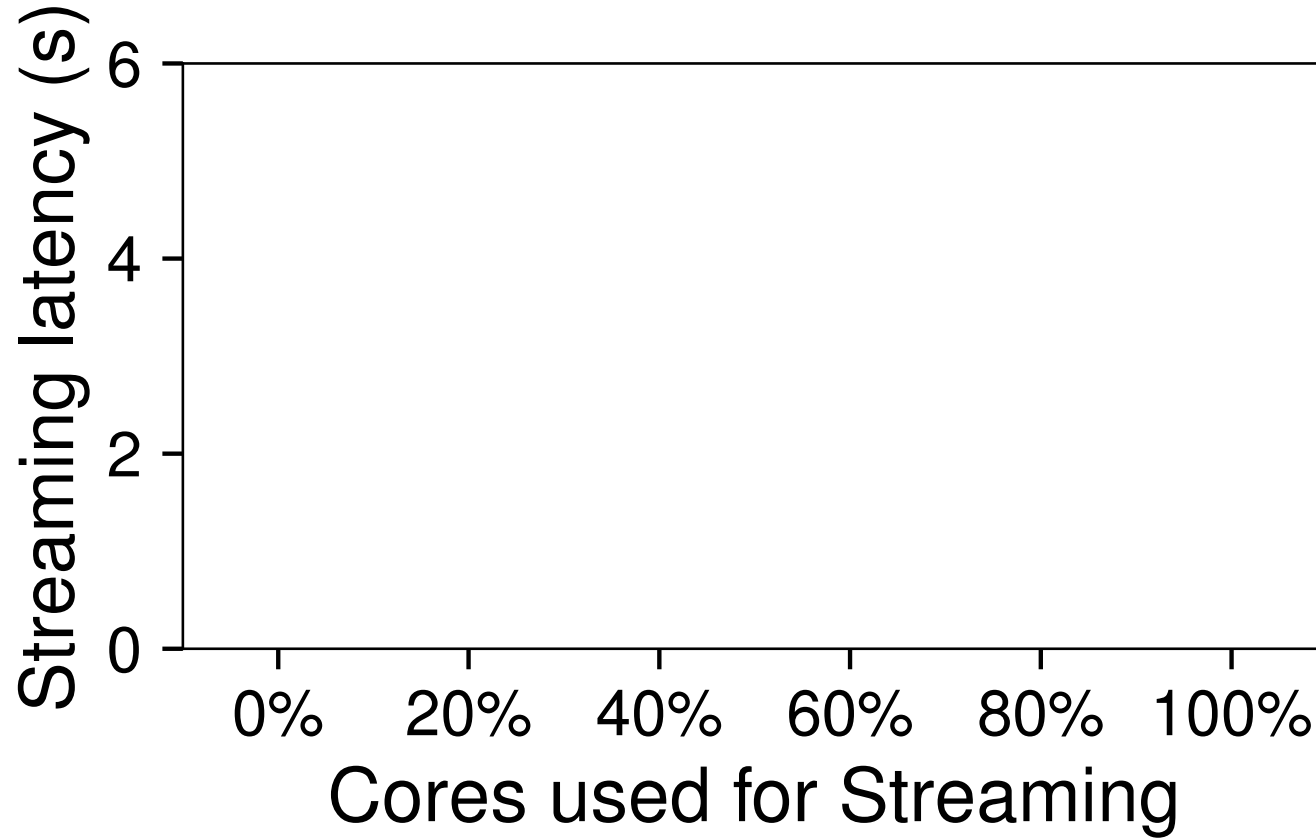
- **LDA**: ML training/inference application uncovering hidden topics from a group of documents
- **Yahoo Streaming Benchmark**: ad-analytics on a stream of ad impressions
- **TPC-H** decision support benchmark

# Benefit of NEPTUNE in stream latency



NEPTUNE achieves latencies comparable to the ideal for the latency-sensitive jobs

# Impact of resource demands in performance



**Past to future**



Efficiently share resources with low impact on throughput

# Summary



**NEPTUNE supports complex unified applications with diverse job requirements!**

- > Suspendable tasks using coroutines
- > Pluggable scheduling policies
- > Continuous unified analytics




<https://github.com/lsds/Neptune>

Thank you!  
Questions?

Panagiotis Garefalakis  
[pgaref@imperial.ac.uk](mailto:pgaref@imperial.ac.uk)

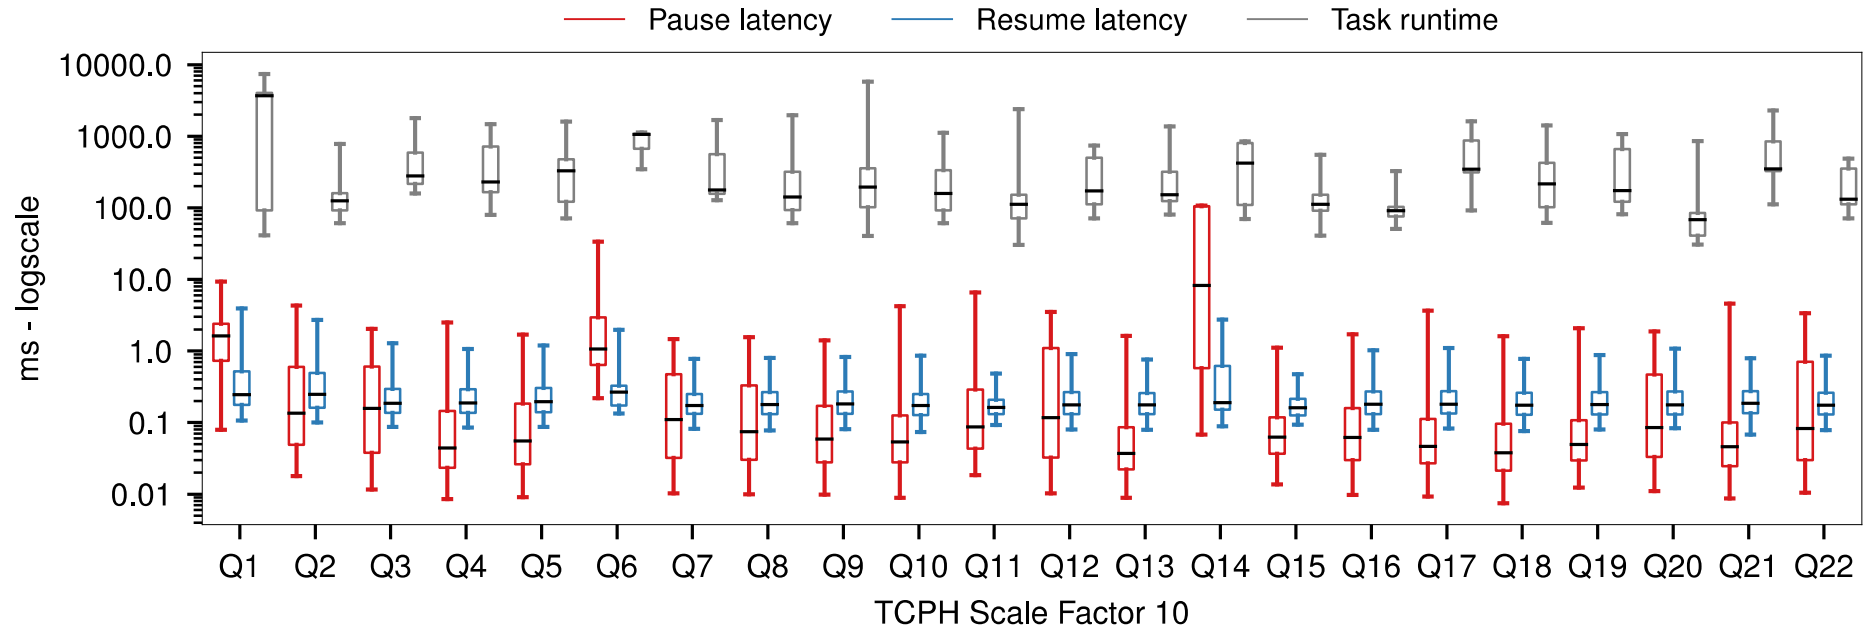
# BACKUP SLIDES

# Suspension mechanism effectiveness



- > **TPCH:** Task runtime distribution for each query ranges from 100s of milliseconds to 10s of seconds

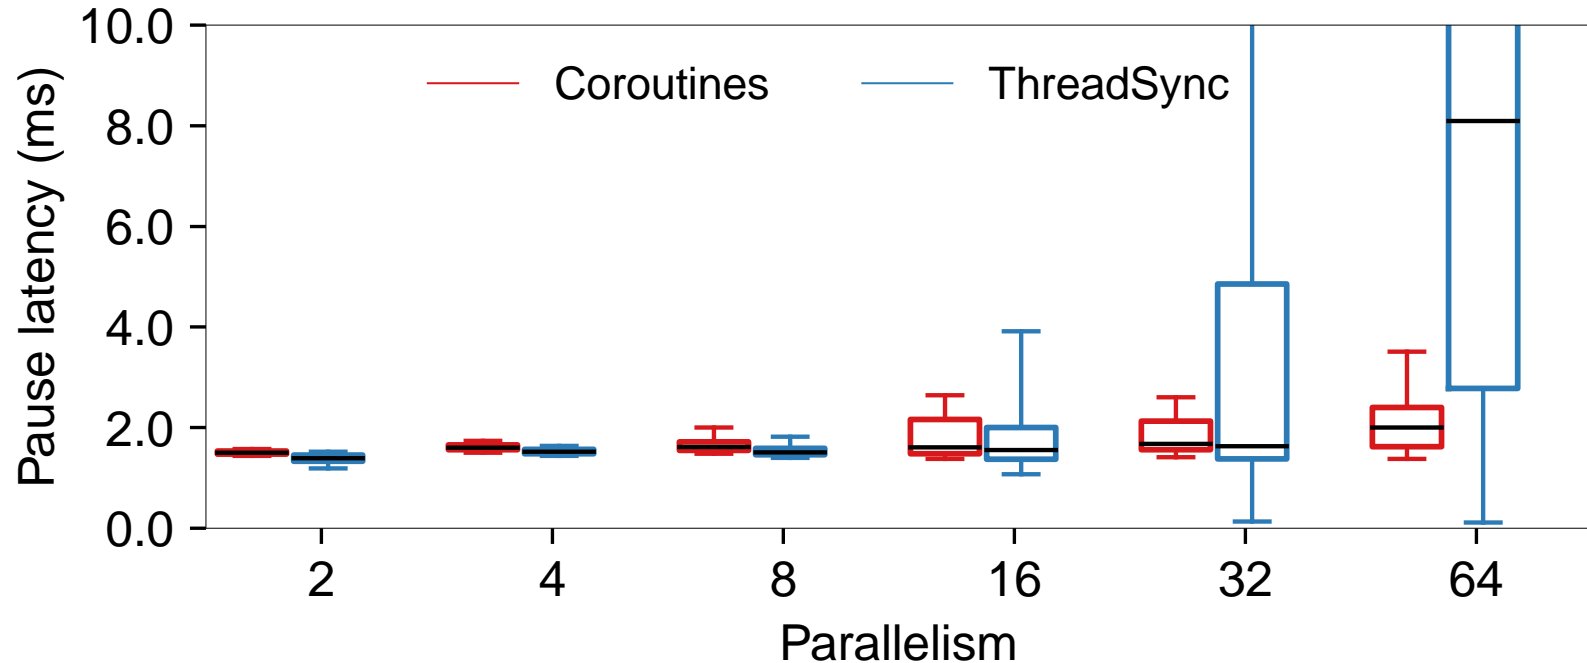
# Suspension mechanism effectiveness



Suspendable tasks effectively pause and resume with sub-millisecond latencies



# Suspension mechanism effectiveness



Coroutine tasks have minimal performance overhead by bypassing the OS

# Demo

---

- > Run a simple **unified** application with
  - > A **high-priority** latency-sensitive job
  - > A **low-priority** latency-tolerant job
  - > Schedule them with default **Spark** and **Neptune**
- > **Goal:** show benefit of Neptune and ease of use

# Suspendable tasks

## Subroutine

```
val collect (TaskContext, Iterator[T]) => (Int, Array[T]) = {  
  (Int, Array[T]) = {  
    val result = new  
    mutable.ArrayBuffer[T]  
    while (itr.hasNext) {  
      result.append(itr.next)  
    }  
    result.toArray  
  }  
}
```

## Coroutine

```
val collect (TaskContext, Iterator[T]) => (Int, Array[T]) = {  
  coroutine {(context: TaskContext, itr: Iterator[T]) => {  
    val result = new mutable.ArrayBuffer[T]  
    while (itr.hasNext) {  
      result.append(itr.next)  
      if (context.isPaused())  
        yieldval(0)  
    }  
    result.toArray  
  }}  
}
```